NORTHSTAR TECHNOLOGIES

OnboardIQ Knowledge Base Catalog

Document inventory, ownership, access scope, and ingestion metadata

Document ID	ONB-CAT-009
Audience	OnboardIQ developers and knowledge administrators
Version	1.0 — Demo
Owner	People Operations / relevant department
Status	Sample content — approval required

Important: Fictional sample company and policy pack created for OnboardIQ demonstration and ingestion testing. These documents are not legal advice and are not actual company policies. Replace with approved company content before real use.

1. Purpose

- Use this catalog to test the document-ingestion pipeline and metadata filtering in OnboardIQ. Each PDF has a stable document ID and a suggested audience.

2. Document inventory

File	Document ID	Suggested audience
01_Employee_Handbook.pdf	HR-HBK-001	All employees
02_New_Joiner_Onboarding_Guide.pdf	HR-ONB-002	New employees and onboarding buddies
03_IT_Acceptable_Use_and_Support_Guide.pdf	IT-USE-003	All employees; technical details for IT staff
04_Information_Security_and_Incident_Response.pdf	SEC-IR-004	All employees; incident handling for Security
05_Leave_Attendance_and_Employee_Support.pdf	HR-SUP-005	All employees
06_Engineering_Developer_Onboarding.pdf	ENG-DEV-006	Engineering and technical teams
07_Data_Privacy_and_Records_Handling.pdf	PRIV-DATA-007	All employees; additional duties for data owners
08_Training_and_Learning_Guide.pdf	LND-TRN-008	All employees and managers

3. Suggested metadata schema

Field	Example	Purpose
document_id	HR-HBK-001	Stable identifier
title	Employee Handbook	Display and retrieval
department	HR	Department filter
access_scope	all_employees	Authorization filter
version	1.0-demo	Version tracking
status	sample_requires_approval	Prevent unapproved content from being treated as official
source_file	01_Employee_Handbook.pdf	Traceability
ingested_at	ISO-8601 timestamp	Audit and refresh tracking

4. Recommended ingestion flow

- 1. Upload PDF and validate type, size, and access permissions.
- 2. Extract text and record page-level provenance.
- 3. Clean repeated headers and footers without removing meaningful policy text.
- 4. Split text into chunks that preserve headings and section context.
- 5. Attach document ID, title, department, access scope, version, and page metadata to each chunk.

- 6. Generate embeddings and upsert vectors into the selected vector database.
- 7. Mark ingestion status and retain an error reason if processing fails.
- 8. At query time, authorize the user first and filter retrieval by allowed access scope before sending context to the LLM.

5. Testing questions

- Where should a new employee report a suspicious email?
- What should an employee do before sharing confidential information?
- What is the suggested process for requesting software access?
- What should an engineer include in a pull request?
- Which team should handle payroll discrepancies?
- What should an employee do if a company device is lost?

6. Important implementation note

- These are synthetic documents. For real deployment, replace them with approved company policies, verify the content with policy owners, assign correct access scopes, and test that users cannot retrieve documents outside their permissions.
- The AI assistant should answer from retrieved, authorized content, cite document titles and pages when possible, and say when the knowledge base does not contain a reliable answer.

End of OnboardIQ Knowledge Base Catalog. This is fictional demonstration content for OnboardIQ.