NORTHSTAR TECHNOLOGIES
PVT. LTD.

AI ACCEPTABLE USE POLICY

Responsible, secure and transparent use of artificial intelligence tools and services

Document ID	IT-POL-003
Document Type	POLICY
Access Scope	ALL_EMPLOYEES
Department	Enterprise-wide
Version	1.0
Effective Date	04 November 2026
Document Owner	IT Operations, with Engineering and HR consultation
Review Cycle	At least annually and after material AI, privacy or security changes

Purpose. This policy sets expectations for the responsible use of generative AI, machine-learning systems, assistants, code-generation tools and automated decision-support services in company work. It aims to enable useful innovation while protecting company information, employees, customers, systems and intellectual property.

Fictional demonstration document for the OnBoardIQ sample knowledge base. This is an illustrative policy, not a formally approved corporate policy, legal advice or a substitute for applicable law, contracts or vendor terms.

1. Scope and applicability

This policy applies to all employees, interns, contractors and approved service providers who use AI-enabled tools or services for Northstar Technologies business. It applies whether the tool is accessed through a browser, desktop application, mobile application, API, IDE extension, internal platform or automated workflow.

The ALL_EMPLOYEES access scope makes this policy available to all authenticated employees. It does not grant approval to use every AI tool or permission to submit every category of company data.

2. Definitions

- **Artificial intelligence (AI):** software that performs tasks such as generating content, classifying information, predicting outcomes, retrieving knowledge or supporting decisions.
- **Generative AI:** AI that creates or transforms text, code, images, audio, video or other content in response to input.
- **AI tool:** an approved or unapproved model, application, API, plugin, extension, agent or hosted service that provides AI capabilities.
- **Company data:** information created, received, processed or maintained for Northstar Technologies business, including internal, confidential and personal information.
- **Human review:** meaningful evaluation by a person with suitable knowledge and authority before an AI output is relied upon for a consequential action.

3. Approved tools and business purposes

- Use only AI tools that have been approved through the applicable company technology, security and procurement process for the intended data classification and use case.
- Use company-managed accounts and configurations where provided; do not route business work through personal accounts to bypass controls.
- Use AI for permitted productivity tasks such as drafting, summarization, brainstorming, coding assistance, knowledge discovery, translation and analysis when the data and use case are approved.
- Confirm whether a tool stores prompts, uses submissions for model improvement, retains files, permits human review or shares data with subprocessors before submitting company information.
- Do not connect an AI service to company repositories, databases, email, storage or internal APIs without authorization and a reviewed integration design.

4. Data classification and information handling

The sensitivity of input data determines whether it may be submitted to an AI service. When uncertain, treat the information as confidential and ask the relevant data owner or security contact before sharing.

- **Public:** may be used with approved tools when no other restriction applies.
- **Internal:** may be used only with tools approved for internal company information and the relevant business purpose.
- **Confidential:** must not be submitted to external AI services unless the service and use case are explicitly approved with suitable contractual, privacy and security controls.
- **Restricted:** must not be submitted to AI tools unless a specific, documented authorization and secure processing arrangement explicitly permits it.
- Never submit passwords, API keys, access tokens, private keys, session cookies or other authentication secrets to an AI prompt.
- Do not upload employee, applicant, customer or partner personal information unless the use case, tool and processing are authorized.
- Remove unnecessary identifiers and minimize the amount of data included in prompts, examples, attachments and context windows.

- Treat embeddings, vector indexes, prompts, generated outputs and conversation histories as potentially sensitive if they contain or reveal company information.

5. Accuracy, verification and human accountability

- AI-generated content may be inaccurate, incomplete, outdated, biased or fabricated. Verify important claims against authoritative sources.
- Review generated code for correctness, security, licensing, maintainability and compatibility before merging or deploying it.
- Validate calculations, data transformations, citations, summaries and recommendations before using them in business decisions.
- Do not represent AI-generated statements as verified facts without appropriate review.
- The employee using the tool remains responsible for the work product and actions taken based on its output.
- Escalate material errors or unsafe outputs that could affect customers, employees, security or business operations.

6. Prohibited uses

- Using AI to facilitate fraud, phishing, impersonation, harassment, discrimination, unauthorized surveillance or other unlawful or harmful activity.
- Attempting to bypass security controls, access restrictions, rate limits or vendor safeguards.
- Submitting confidential or restricted data to an unapproved tool or personal AI account.
- Using AI output as the sole basis for hiring, promotion, termination, compensation, disciplinary action or another consequential employment decision.
- Generating or distributing deceptive content that falsely claims to be an official company communication or an authentic statement by a person.
- Using AI-generated code, scripts or agents to access systems or data without authorization.
- Allowing an AI agent to execute production changes, send external communications, approve transactions or alter access rights without an approved control design and required human authorization.
- Removing provenance, attribution or required disclosures to mislead recipients about AI-generated material.

7. AI-assisted software development

- Use only approved coding assistants and integrations for company source code and technical documentation.
- Do not submit secrets, production data, customer records or proprietary code to tools that are not approved for that information.
- Treat generated code as untrusted third-party input: inspect dependencies, licenses, security implications and hidden behaviours.
- Run tests, static analysis and applicable security scans before accepting AI-assisted changes.
- Do not paste generated commands into production terminals without understanding their effect and following change controls.
- Ensure code contributions meet repository review, testing and release requirements; AI assistance does not replace engineering accountability.

8. AI agents, plugins and automated actions

- An agent must have a defined owner, approved purpose, documented permissions and bounded operating scope.
- Grant the agent only the minimum tools, data sources and permissions required for its task.
- Require human approval for consequential or externally visible actions unless a documented risk assessment and approved control framework explicitly permit automation.

- Use allowlists, input validation, output validation, execution limits and audit logs for tool-enabled agents.
- Protect against prompt injection and malicious instructions embedded in retrieved documents, web pages, emails, code comments or uploaded files.
- Separate read-only retrieval from write-capable tools and use explicit authorization checks at the service boundary.
- Provide a way to stop, revoke or disable an agent and investigate its actions.

9. Privacy, intellectual property and attribution

- Use personal information only for an approved purpose and in accordance with applicable privacy requirements and company procedures.
- Do not assume that generated content is original, accurate or free of third-party rights. Review sources, licenses and terms before commercial use.
- Respect confidentiality, copyright, open-source licenses, contractual restrictions and customer ownership requirements.
- Do not upload third-party confidential information unless the company is authorized to share it with the selected service.
- Follow project or customer requirements for disclosure of AI assistance, attribution and provenance.
- Retain prompts and generated artifacts only as long as required for the approved business purpose and applicable retention rules.

10. Transparency and external communications

- Review AI-generated external communications before sending them on behalf of the company.
- Do not fabricate testimonials, customer approvals, research results, employee statements or evidence.
- Disclose AI involvement when required by law, contract, customer commitment or applicable company procedure.
- Ensure a human owner is identifiable for important decisions and communications supported by AI.

11. Reporting security, privacy or quality concerns

Report suspected disclosure of company information, exposure of credentials, unsafe agent actions, unauthorized integrations, harmful outputs or material inaccuracies through the approved security, privacy or incident-reporting channel.

Include the tool name, approximate time, affected data category, relevant account or workflow and steps already taken. Do not forward sensitive content more widely than necessary.

If a secret may have been exposed, contact the responsible security or system owner promptly so credentials can be revoked or rotated. Preserve relevant evidence and follow incident-response instructions.

12. Training and awareness

- Complete required AI, information-security and privacy awareness training.
- Understand the approved-tool list and data-classification rules before using AI for company work.
- Managers should ensure team members understand relevant risks for their role and use case.
- Teams building or deploying AI systems should receive additional training appropriate to model risk, data handling, evaluation and operational controls.

13. Exceptions and approvals

A request to use an unapproved tool, submit a higher-sensitivity data category or enable an agent with elevated permissions must be reviewed before use. The request should identify the business purpose, data involved, service provider, retention and training settings, access model, risk assessment and proposed safeguards.

Approval must come from the appropriate data owner and designated security, privacy, legal or technology authority under the company's governance process. An exception must have a documented scope, owner, compensating controls

and expiry or review date.

14. Employee quick checklist

- Is the AI tool approved for this business purpose?
- Am I allowed to share this data with this tool?
- Have I removed secrets and unnecessary personal or confidential information?
- Will I verify the output before relying on it?
- Could this output affect a person's rights, employment, finances, safety or access?
- Does an agent have only the permissions it needs, with approval gates for consequential actions?
- Do I need to disclose AI assistance or preserve provenance?
- Do I know how to report a suspected issue?

15. Policy ownership and review

IT Operations coordinates policy maintenance with Engineering, HR, Security, Privacy and other relevant stakeholders. The policy should be reviewed at least annually and whenever there are material changes to AI services, company data classifications, legal requirements, security threats or business use cases.

Violations may result in removal of tool access, required remediation, escalation under applicable company procedures or other action consistent with law and contractual obligations.

Document governance

Policy owner: IT Operations. Effective date: 04 November 2026. Version: 1.0. Changes should be version-controlled and communicated to all affected employees.

End of document. Fictional Northstar Technologies Pvt. Ltd. demonstration artifact for the OnBoardIQ sample knowledge base. All procedures and examples are illustrative.